

ad-click-prediction

September 11, 2024

1 Understanding the Problem

[18]:

2 Load and Explore the Data

```
[19]: import pandas as pd
import seaborn as sns

df = pd.read_csv('/content/ad_click_dataset.csv')
df.set_index('id', inplace = True)
df.head()
```

```
[19]:      full_name  age      gender device_type ad_position browsing_history \
id
670   User670   22.0      NaN    Desktop      Top      Shopping
3044  User3044   NaN      Male    Desktop      Top      NaN
5912  User5912  41.0  Non-Binary      NaN      Side      Education
5418  User5418  34.0      Male      NaN      NaN      Entertainment
9452  User9452  39.0  Non-Binary      NaN      NaN      Social Media

      time_of_day  click
id
670   Afternoon      1
3044      NaN      1
5912    Night      1
5418   Evening      1
9452   Morning      0
```

```
[20]: df.shape
```

```
[20]: (10000, 8)
```

```
[21]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 10000 entries, 670 to 3056
```

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	full_name	10000 non-null	object
1	age	5234 non-null	float64
2	gender	5307 non-null	object
3	device_type	8000 non-null	object
4	ad_position	8000 non-null	object
5	browsing_history	5218 non-null	object
6	time_of_day	8000 non-null	object
7	click	10000 non-null	int64

dtypes: float64(1), int64(1), object(6)

memory usage: 703.1+ KB

```
[22]: df.describe()
```


```
[22]:
```

	age	click
count	5234.000000	10000.000000
mean	40.197363	0.650000
std	13.126420	0.476993
min	18.000000	0.000000
25%	29.000000	0.000000
50%	39.500000	1.000000
75%	52.000000	1.000000
max	64.000000	1.000000

```
[23]: df.columns
```

```
[23]: Index(['full_name', 'age', 'gender', 'device_type', 'ad_position',  
        'browsing_history', 'time_of_day', 'click'],  
        dtype='object')
```


```
[24]: import matplotlib.pyplot as plt  
df.hist(bins=30, figsize=(15, 10))  
plt.tight_layout()  
plt.show()
```



```
[25]: corr_matrix = df.select_dtypes([int, float]).corr()
      print(corr_matrix)
```

```
          age    click
age    1.000000 -0.082056
click -0.082056  1.000000
```

```
[26]: plt.figure(figsize=(12, 8))
      sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
      plt.show()
```



3 Data Preprocessing

3.1 Encode Categorical Variables

```
[27]: from sklearn.preprocessing import LabelEncoder
obj_cols = df.select_dtypes(object).columns
for col in obj_cols:
    encoder = LabelEncoder()
    df[col] = encoder.fit_transform(df[col])

df.dtypes
```

```
[27]: full_name      int64
age                float64
gender             int64
device_type        int64
ad_position         int64
browsing_history    int64
time_of_day         int64
click              int64
```

dtype: object

3.2 Handling Missing Values

```
[28]: df['age'].fillna(df['age'].mean(), inplace=True)
```

```
[29]: df.isnull().sum()
```

```
[29]: full_name      0
      age           0
      gender        0
      device_type    0
      ad_position    0
      browsing_history 0
      time_of_day    0
      click          0
      dtype: int64
```

```
[30]: import pandas as pd

      duplicates = df.duplicated()

      duplicates_in_subset = df.duplicated(subset=['full_name', 'age', 'gender',
      ↪ 'device_type', 'ad_position', 'browsing_history', 'time_of_day', 'click'])

      # Remove all duplicate rows based on all columns
      df_cleaned = df.drop_duplicates()

      # Remove duplicates based on a subset of columns
      df_cleaned = df.drop_duplicates(subset=['full_name', 'age', 'gender',
      ↪ 'device_type', 'ad_position', 'browsing_history', 'time_of_day', 'click'])

      df_cleaned = df.groupby(['full_name', 'age', 'gender', 'device_type',
      ↪ 'ad_position', 'browsing_history', 'time_of_day']).agg({'click': 'sum'}).
      ↪reset_index()

      print(df_cleaned.info())
      print(df_cleaned.duplicated().sum()) # Should return 0 if all duplicates are
      ↪removed
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7147 entries, 0 to 7146
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   full_name       7147 non-null  int64
1   age             7147 non-null  float64
```

```

2   gender          7147 non-null   int64
3   device_type     7147 non-null   int64
4   ad_position     7147 non-null   int64
5   browsing_history 7147 non-null   int64
6   time_of_day     7147 non-null   int64
7   click           7147 non-null   int64
dtypes: float64(1), int64(7)
memory usage: 446.8 KB
None
0

```

4 Split the Data

```
[31]: x = df.drop(columns = ['click'])
      y = df['click']
```

```
[32]: from sklearn.model_selection import train_test_split

      xtrain, xtest, ytrain, ytest = train_test_split(x, y, train_size = 0.8)
```

5 Model Selection

```
[33]: from sklearn.metrics import accuracy_score

      def eval_model(model, xtrain, ytrain, xtest, ytest):
          model.fit(xtrain, ytrain)
          trainpred = model.predict(xtrain)
          testpred = model.predict(xtest)
          return accuracy_score(ytrain, trainpred), accuracy_score(ytest, testpred)
```

```
[34]: result = pd.DataFrame(columns = ['Model Number', 'Name', 'Training Accuracy',
    ↪ 'Testing Accuracy'])
      result.head()
```

```
[34]: Empty DataFrame
      Columns: [Model Number, Name, Training Accuracy, Testing Accuracy]
      Index: []
```

```
[35]: from sklearn.tree import DecisionTreeClassifier
      model = DecisionTreeClassifier()
      row = []
      row.extend([1, 'Decision Tree'])
      row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
      result.loc[len(result.index)] = row
      result.head()
```

```
[35]:
```

Model Number	Name	Training Accuracy	Testing Accuracy
0	1 Decision Tree	1.0	0.8325

```
[36]: from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier()
row = []
row.extend([2, 'KNN'])
row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
result.loc[len(result.index)] = row
result.head()
```

```
[36]:
```

Model Number	Name	Training Accuracy	Testing Accuracy
0	1 Decision Tree	1.000000	0.8325
1	2 KNN	0.865125	0.8230

```
[37]: from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
row = []
row.extend([3, 'Extra Trees'])
row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
result.loc[len(result.index)] = row
result.head()
```

```
[37]:
```

Model Number	Name	Training Accuracy	Testing Accuracy
0	1 Decision Tree	1.000000	0.8325
1	2 KNN	0.865125	0.8230
2	3 Extra Trees	1.000000	0.7720

```
[38]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
row = []
row.extend([4, 'Logistic Regression'])
row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
result.loc[len(result.index)] = row
result.head()
```

```
[38]:
```

Model Number	Name	Training Accuracy	Testing Accuracy
0	1 Decision Tree	1.000000	0.8325
1	2 KNN	0.865125	0.8230
2	3 Extra Trees	1.000000	0.7720
3	4 Logistic Regression	0.644875	0.6705

```
[39]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
row = []
row.extend([5, 'Random Forest'])
row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
```

```
result.loc[len(result.index)] = row
result.head()
```

```
[39]:
```

	Model Number	Name	Training Accuracy	Testing Accuracy
0	1	Decision Tree	1.000000	0.8325
1	2	KNN	0.865125	0.8230
2	3	Extra Trees	1.000000	0.7720
3	4	Logistic Regression	0.644875	0.6705
4	5	Random Forest	1.000000	0.7990

```
[40]: from sklearn.svm import SVC
model = SVC()
row = []
row.extend([6, 'SVM'])
row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
result.loc[len(result.index)] = row
result.head()
```

```
[40]:
```

	Model Number	Name	Training Accuracy	Testing Accuracy
0	1	Decision Tree	1.000000	0.8325
1	2	KNN	0.865125	0.8230
2	3	Extra Trees	1.000000	0.7720
3	4	Logistic Regression	0.644875	0.6705
4	5	Random Forest	1.000000	0.7990

```
[41]: from sklearn.ensemble import GradientBoostingClassifier
model = GradientBoostingClassifier()
row = []
row.extend([7, 'Gradient Boosting'])
row.extend(eval_model(model, xtrain, ytrain, xtest, ytest))
result.loc[len(result.index)] = row
result
```

```
[41]:
```

	Model Number	Name	Training Accuracy	Testing Accuracy
0	1	Decision Tree	1.000000	0.8325
1	2	KNN	0.865125	0.8230
2	3	Extra Trees	1.000000	0.7720
3	4	Logistic Regression	0.644875	0.6705
4	5	Random Forest	1.000000	0.7990
5	6	SVM	0.644875	0.6705
6	7	Gradient Boosting	0.712625	0.7160

```
[42]: import numpy as np

models = result['Name']
training_accuracies = np.round(result['Training Accuracy'],7)
testing_accuracies = np.round(result['Testing Accuracy'],7)
```



```

x = np.arange(len(models)) # the label locations
width = 0.35 # the width of the bars

fig, ax = plt.subplots(figsize=(12, 8))
rects1 = ax.bar(x - width/2, training_accuracies, width, label='Training_
↳Accuracy')
rects2 = ax.bar(x + width/2, testing_accuracies, width, label='Testing_
↳Accuracy')

# Add some text for labels, title and custom x-axis tick labels, etc.
ax.set_xlabel('Model')
ax.set_ylabel('Accuracy')
ax.set_title('Training and Testing Accuracy of Various ML Models')
ax.set_xticks(x)
ax.set_xticklabels(models)
ax.legend()

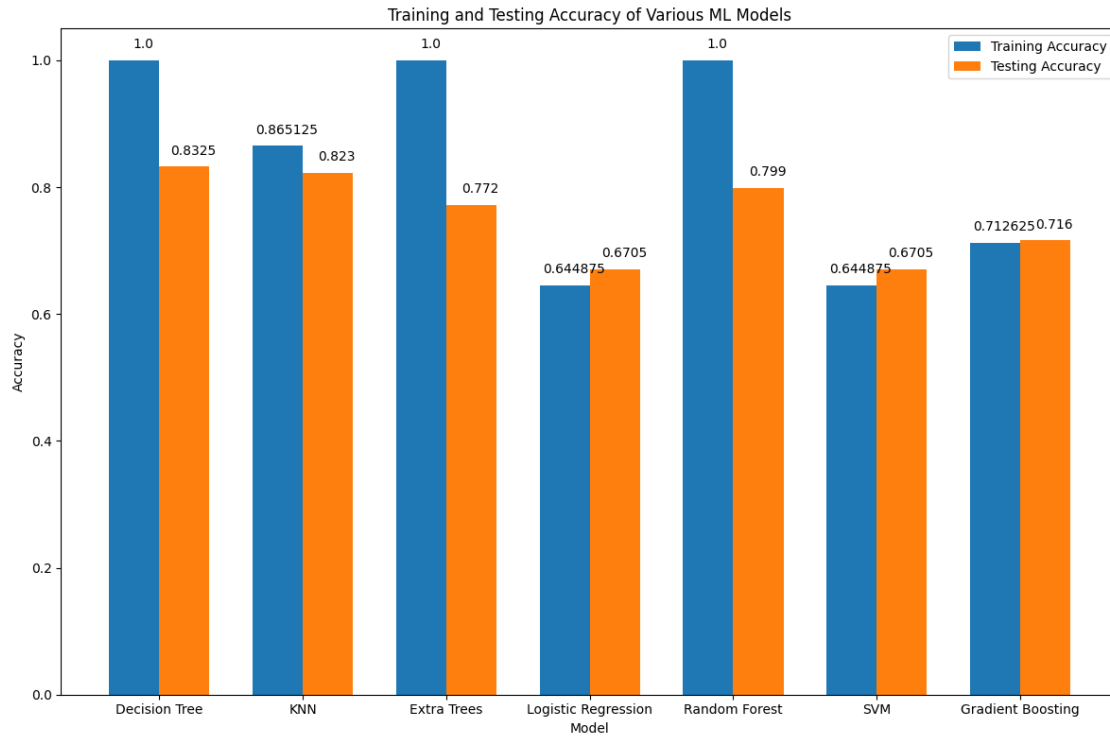
# Function to add labels on top of the bars
def autolabel(rects):
    for rect in rects:
        height = rect.get_height()
        ax.annotate('{}' .format(height),
                    xy=(rect.get_x() + rect.get_width() / 2, height),
                    xytext=(7, 7), # 7 points vertical offset
                    textcoords="offset points",
                    ha='center', va='bottom')

autolabel(rects1)
autolabel(rects2)

fig.tight_layout()

plt.show()

```



6 Model Training

```
[43]: from sklearn.neighbors import KNeighborsClassifier
      model = KNeighborsClassifier()
      model.fit(xtrain, ytrain)
```

```
[43]: KNeighborsClassifier()
```

7 Model Evaluation

```
[44]: trainpred = model.predict(xtrain)
      testpred = model.predict(xtest)
```

```
[45]: from sklearn.metrics import classification_report
```

```
[46]: print(classification_report(ytrain, trainpred))
```

	precision	recall	f1-score	support
0	0.95	0.65	0.77	2841
1	0.84	0.98	0.90	5159

accuracy			0.87	8000
macro avg	0.90	0.82	0.84	8000
weighted avg	0.88	0.87	0.86	8000

```
[47]: print(classification_report(ytest, testpred))
```

	precision	recall	f1-score	support
0	0.89	0.53	0.66	659
1	0.81	0.97	0.88	1341

accuracy			0.82	2000
macro avg	0.85	0.75	0.77	2000
weighted avg	0.83	0.82	0.81	2000